

#Caricature Image To Real Image

```
from google.colab import drive
drive.mount('/content/drive')
```

```
# extracting the compressed Dataset
from zipfile import ZipFile
dataset = '/content/drive/MyDrive/archive.zip'
from zipfile import ZipFile
```

```
# with ZipFile(dataset, 'r') as zip:
#   zip.extractall()
!unzip '/content/drive/MyDrive/archive.zip' -d '/content/drive/MyDrive'
print('The dataset is extracted')
```



Show hidden output

```
!pip install torch torchvision matplotlib opencv-python tqdm
```



Show hidden output

#Load the dataset

```
# Replace these paths with your real and caricature folder paths
caric_path = "/content/drive/MyDrive/CaVI_Dataset/CaVI_Dataset/Caricature"
real_path = "/content/drive/MyDrive/CaVI_Dataset/CaVI_Dataset/Real"
```

```
import cv2
import os
```

```
def crop_faces(input_folder, output_folder):
    face_cascade = cv2.CascadeClassifier(cv2.data.haarcascades + 'haarcascade_frontalface_default.xml')
    os.makedirs(output_folder, exist_ok=True)
    for img_name in os.listdir(input_folder):
        img_path = os.path.join(input_folder, img_name)
        img = cv2.imread(img_path)
        gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
        faces = face_cascade.detectMultiScale(gray, 1.3, 5)
        for (x, y, w, h) in faces:
            face = img[y:y+h, x:x+w]
            resized = cv2.resize(face, (256, 256))
            cv2.imwrite(os.path.join(output_folder, img_name), resized)
            break
```

```
import torch.nn as nn
```

```
class ResidualBlock(nn.Module):
    def __init__(self, dim):
        super().__init__()
        self.block = nn.Sequential(
            nn.Conv2d(dim, dim, kernel_size=3, padding=1),
            nn.InstanceNorm2d(dim),
            nn.ReLU(inplace=True),
            nn.Conv2d(dim, dim, kernel_size=3, padding=1),
            nn.InstanceNorm2d(dim),
        )

    def forward(self, x):
        return x + self.block(x)
```

```

class Generator(nn.Module):
    def __init__(self, in_channels=3, out_channels=3, n_res_blocks=6):
        super().__init__()
        model = [
            nn.Conv2d(in_channels, 64, kernel_size=7, stride=1, padding=3),
            nn.InstanceNorm2d(64),
            nn.ReLU(inplace=True)
        ]

        # Downsampling
        model += [
            nn.Conv2d(64, 128, kernel_size=4, stride=2, padding=1),
            nn.InstanceNorm2d(128),
            nn.ReLU(inplace=True),
            nn.Conv2d(128, 256, kernel_size=4, stride=2, padding=1),
            nn.InstanceNorm2d(256),
            nn.ReLU(inplace=True)
        ]

        # Residual blocks
        for _ in range(n_res_blocks):
            model += [ResidualBlock(256)]

        # Upsampling
        model += [
            nn.ConvTranspose2d(256, 128, kernel_size=4, stride=2, padding=1),
            nn.InstanceNorm2d(128),
            nn.ReLU(inplace=True),
            nn.ConvTranspose2d(128, 64, kernel_size=4, stride=2, padding=1),
            nn.InstanceNorm2d(64),
            nn.ReLU(inplace=True),
            nn.Conv2d(64, out_channels, kernel_size=7, stride=1, padding=3),
            nn.Tanh()
        ]

        self.model = nn.Sequential(*model)

    def forward(self, x):
        return self.model(x)

class Discriminator(nn.Module):
    def __init__(self, in_channels=3):
        super().__init__()
        def discriminator_block(in_filters, out_filters, normalize=True):
            layers = [nn.Conv2d(in_filters, out_filters, kernel_size=4, stride=2, padding=1)]
            if normalize:
                layers.append(nn.InstanceNorm2d(out_filters))
            layers.append(nn.LeakyReLU(0.2, inplace=True))
            return layers

        self.model = nn.Sequential(
            *discriminator_block(in_channels, 64, normalize=False),
            *discriminator_block(64, 128),
            *discriminator_block(128, 256),
            *discriminator_block(256, 512),
            nn.Conv2d(512, 1, kernel_size=4, padding=1)
        )

    def forward(self, img):
        return self.model(img)

from torch.utils.data import Dataset, DataLoader
from PIL import Image

```

```

import glob
import torchvision.transforms as transforms
import os

class FaceCaricatureDataset(Dataset):
    def __init__(self, real_root, caric_root, transform=None):
        # Modified to include various image extensions and ensure paths are correct
        self.real_paths = sorted(glob.glob(os.path.join(real_root, '**/*.jpg'), recursive=True) +
                                glob.glob(os.path.join(real_root, '**/*.png'), recursive=True) +
                                glob.glob(os.path.join(real_root, '**/*.jpeg'), recursive=True))
        self.caric_paths = sorted(glob.glob(os.path.join(caric_root, '**/*.jpg'), recursive=True) +
                                glob.glob(os.path.join(caric_root, '**/*.png'), recursive=True) +
                                glob.glob(os.path.join(caric_root, '**/*.jpeg'), recursive=True))
        self.transform = transform

    def __len__(self):
        return min(len(self.real_paths), len(self.caric_paths)) # Ensure both paths have same number of images

    def __getitem__(self, index):
        real_img = Image.open(self.real_paths[index]).convert('RGB')
        caric_img = Image.open(self.caric_paths[index]).convert('RGB')
        if self.transform:
            real_img = self.transform(real_img)
            caric_img = self.transform(caric_img)
        return real_img, caric_img

transform = transforms.Compose([
    transforms.Resize((256, 256)),
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])

dataset = FaceCaricatureDataset(caric_path, real_path, transform=transform)
print(len(dataset)) # Print dataset length to check for 0
print(dataset.real_paths[:5]) # Print a few real image paths for verification
print(dataset.caric_paths[:5]) # Print a few caricature image paths for verification
dataloader = DataLoader(dataset, batch_size=4, shuffle=True)

```

4816

```

['/content/drive/MyDrive/CaVI_Dataset/CaVI_Dataset/Caricature/Aamir_Khan/Aamir_Khan_c_0.jpg', '/content/drive/MyDr
 ['/content/drive/MyDrive/CaVI_Dataset/CaVI_Dataset/Real/Aamir_Khan/Aamir_Khan_r_0.jpg', '/content/drive/MyDrive/Ca

```

#Train the Model

```

import torch
import torch.optim as optim
import torch.nn as nn
from tqdm import tqdm

# Check for GPU availability and fallback to CPU if not available
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device) # Inform the user of the device being used

G = Generator().to(device)
D = Discriminator().to(device)

optimizer_G = optim.Adam(G.parameters(), lr=0.0002, betas=(0.5, 0.999))
optimizer_D = optim.Adam(D.parameters(), lr=0.0002, betas=(0.5, 0.999))

adversarial_loss = nn.MSELoss()
reconstruction_loss = nn.L1Loss()

for epoch in range(5):

```

```

for i, (real, caric) in enumerate(tqdm(data_loader)):
    # Move data to the selected device
    real, caric = real.to(device), caric.to(device)
    valid = torch.ones((real.size(0), 1, 15, 15)).to(device)
    fake = torch.zeros((real.size(0), 1, 15, 15)).to(device)

    # -----
    # Train Generator
    # -----
    optimizer_G.zero_grad()
    fake_caric = G(real)
    g_adv = adversarial_loss(D(fake_caric), valid)
    g_recon = reconstruction_loss(fake_caric, caric)
    g_loss = g_adv + 10 * g_recon
    g_loss.backward()
    optimizer_G.step()

    # -----
    # Train Discriminator
    # -----
    optimizer_D.zero_grad()
    real_loss = adversarial_loss(D(caric), valid)
    fake_loss = adversarial_loss(D(fake_caric.detach()), fake)
    d_loss = 0.5 * (real_loss + fake_loss)
    d_loss.backward()
    optimizer_D.step()

print(f"Epoch {epoch} | Generator Loss: {g_loss.item():.4f} | Discriminator Loss: {d_loss.item():.4f}")

torch.save({
    'epoch': epoch,
    'generator': G.state_dict(),
    'discriminator': D.state_dict(),
    'optimizer_G': optimizer_G.state_dict(),
    'optimizer_D': optimizer_D.state_dict()
}, f'cafe_gan_epoch_{epoch+1}.pth')

torch.save({'generator': G.state_dict(), 'discriminator': D.state_dict()}, 'cafe_gan_models.pth')

```

```

➡ Using device: cuda
100%|██████████| 1204/1204 [05:54<00:00, 3.39it/s]
Epoch 0 | Generator Loss: 5.8244 | Discriminator Loss: 0.1455
100%|██████████| 1204/1204 [05:54<00:00, 3.40it/s]
Epoch 1 | Generator Loss: 5.9509 | Discriminator Loss: 0.1566
100%|██████████| 1204/1204 [05:54<00:00, 3.39it/s]
Epoch 2 | Generator Loss: 5.2251 | Discriminator Loss: 0.1412
100%|██████████| 1204/1204 [05:54<00:00, 3.39it/s]
Epoch 3 | Generator Loss: 5.5049 | Discriminator Loss: 0.1549
100%|██████████| 1204/1204 [05:54<00:00, 3.39it/s]
Epoch 4 | Generator Loss: 6.5772 | Discriminator Loss: 0.1155

```

```

import os
from torchvision.utils import save_image

os.makedirs("output", exist_ok=True)

G.eval()
with torch.no_grad():
    for i, (real, _) in enumerate(data_loader):
        real = real.cuda()
        fake_caric = G(real)
        save_image(fake_caric, f'output/caric_{i}.png', normalize=True)
        save_image(real, f'output/real_{i}.png', normalize=True)
        if i == 2: break

```

```
# Save
torch.save(G.state_dict(), "/content/drive/MyDrive/generator.pth") # Corrected path
torch.save(D.state_dict(), "/content/drive/MyDrive/discriminator.pth") # Corrected path
```

#Load the Model

```
!pip install torch torchvision matplotlib opencv-python tqdm
```

```
Requirement already satisfied: torch in /usr/local/lib/python3.11/dist-packages (2.6.0+cu124)
Requirement already satisfied: torchvision in /usr/local/lib/python3.11/dist-packages (0.21.0+cu124)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (3.10.0)
Requirement already satisfied: opencv-python in /usr/local/lib/python3.11/dist-packages (4.11.0.86)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (4.67.1)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from torch) (3.18.0)
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.11/dist-packages (from torch) (
Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages (from torch) (3.4.2)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.11/dist-packages (from torch) (3.1.6)
Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from torch) (2025.3.2)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from t
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from t
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from t
Requirement already satisfied: nvidia-cudnn-cu12==9.1.0.70 in /usr/local/lib/python3.11/dist-packages (from torch)
Requirement already satisfied: nvidia-cublas-cu12==12.4.5.8 in /usr/local/lib/python3.11/dist-packages (from torch
Requirement already satisfied: nvidia-cufft-cu12==11.2.1.3 in /usr/local/lib/python3.11/dist-packages (from torch)
Requirement already satisfied: nvidia-curand-cu12==10.3.5.147 in /usr/local/lib/python3.11/dist-packages (from tor
Requirement already satisfied: nvidia-cusolver-cu12==11.6.1.9 in /usr/local/lib/python3.11/dist-packages (from tor
Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in /usr/local/lib/python3.11/dist-packages (from torc
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /usr/local/lib/python3.11/dist-packages (from torch) (2
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from to
Requirement already satisfied: triton==3.2.0 in /usr/local/lib/python3.11/dist-packages (from torch) (3.2.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-packages (from torch) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from torchvision) (2.0.2)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.11/dist-packages (from torchvision)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (4.5
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.4
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (24.2)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (3.2
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.7->mat
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-packages (from Jinja2->torch) (3.
```

```
import torch.nn as nn
import torch
import torch.optim as optim
import torch.nn as nn
from tqdm import tqdm
from torch.utils.data import Dataset, DataLoader
from PIL import Image
import glob
import torchvision.transforms as transforms
import os
import cv2
import os

# Define the Generator and Discriminator classes first
class ResidualBlock(nn.Module):
    def __init__(self, dim):
        super().__init__()
        self.block = nn.Sequential(
```

```

        nn.Conv2d(dim, dim, kernel_size=3, padding=1),
        nn.InstanceNorm2d(dim),
        nn.ReLU(inplace=True),
        nn.Conv2d(dim, dim, kernel_size=3, padding=1),
        nn.InstanceNorm2d(dim),
    )

    def forward(self, x):
        return x + self.block(x)

class Generator(nn.Module):
    def __init__(self, in_channels=3, out_channels=3, n_res_blocks=6):
        super().__init__()
        model = [
            nn.Conv2d(in_channels, 64, kernel_size=7, stride=1, padding=3),
            nn.InstanceNorm2d(64),
            nn.ReLU(inplace=True)
        ]

        # Downsampling
        model += [
            nn.Conv2d(64, 128, kernel_size=4, stride=2, padding=1),
            nn.InstanceNorm2d(128),
            nn.ReLU(inplace=True),
            nn.Conv2d(128, 256, kernel_size=4, stride=2, padding=1),
            nn.InstanceNorm2d(256),
            nn.ReLU(inplace=True)
        ]

        # Residual blocks
        for _ in range(n_res_blocks):
            model += [ResidualBlock(256)]

        # Upsampling
        model += [
            nn.ConvTranspose2d(256, 128, kernel_size=4, stride=2, padding=1),
            nn.InstanceNorm2d(128),
            nn.ReLU(inplace=True),
            nn.ConvTranspose2d(128, 64, kernel_size=4, stride=2, padding=1),
            nn.InstanceNorm2d(64),
            nn.ReLU(inplace=True),
            nn.Conv2d(64, out_channels, kernel_size=7, stride=1, padding=3),
            nn.Tanh()
        ]

        self.model = nn.Sequential(*model)

    def forward(self, x):
        return self.model(x)

class Discriminator(nn.Module):
    def __init__(self, in_channels=3):
        super().__init__()
        def discriminator_block(in_filters, out_filters, normalize=True):
            layers = [nn.Conv2d(in_filters, out_filters, kernel_size=4, stride=2, padding=1)]
            if normalize:
                layers.append(nn.InstanceNorm2d(out_filters))
            layers.append(nn.LeakyReLU(0.2, inplace=True))
            return layers

        self.model = nn.Sequential(
            *discriminator_block(in_channels, 64, normalize=False),
            *discriminator_block(64, 128),
            *discriminator_block(128, 256),
            *discriminator_block(256, 512),
            nn.Conv2d(512, 1, kernel_size=4, padding=1)

```

```

    )

    def forward(self, img):
        return self.model(img)

# Create instances of Generator and Discriminator
G = Generator()
D = Discriminator()

# Load the state dictionaries with map_location to CPU
G.load_state_dict(torch.load("/content/drive/MyDrive/generator.pth", map_location=torch.device('cpu'))) # Load Generator
D.load_state_dict(torch.load("/content/drive/MyDrive/discriminator.pth", map_location=torch.device('cpu'))) # Load Discriminator

<All keys matched successfully>

```

#Test the Model

```

import os
from torchvision.utils import save_image
from PIL import Image

# Set this path to a folder of test face images (either real or caricature) or a single image
test_path = "/content/drive/MyDrive/CaVI_Dataset/CaVI_Dataset/Caricature/Bill_Gates/Bill_Gates_c_0.jpg"

# Output folder
output_path = "/content/output"
os.makedirs(output_path, exist_ok=True)

# Transform for test images
transform = transforms.Compose([
    transforms.Resize((256, 256)),
    transforms.ToTensor(),
    transforms.Normalize((0.5, ), (0.5, ))
])

# Inference
# Check if CUDA is available, otherwise use CPU
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
G.to(device).eval() # Move the model to the correct device and set to evaluation mode

with torch.no_grad():
    # If test_path is a directory, loop through images
    if os.path.isdir(test_path):
        for img_name in os.listdir(test_path):
            img_path = os.path.join(test_path, img_name)
            image = Image.open(img_path).convert("RGB")
            input_tensor = transform(image).unsqueeze(0).to(device) # Move input tensor to the correct device
            output = G(input_tensor)
            save_path = os.path.join(output_path, f"out_{img_name}")
            save_image(output, save_path, normalize=True)
    # If test_path is a single image file, process it directly
    else:
        image = Image.open(test_path).convert("RGB")
        input_tensor = transform(image).unsqueeze(0).to(device) # Move input tensor to the correct device
        output = G(input_tensor)
        # Get the filename from the path
        img_name = os.path.basename(test_path)
        save_path = os.path.join(output_path, f"out_{img_name}")
        save_image(output, save_path, normalize=True)

print("output: Caricature to Real image \n\n")
from IPython.display import Image
orig_img = Image(filename=test_path)
display(orig_img)

```

```
from PIL import Image  
pil_img = Image(filename=save_path)  
display(pil_img)
```

↔ output: Caricature to Real image

